

Tools used

1. Terraform is used to deploy the EKS cluster and the flask application
2. GitHub is used as the code repository.
3. GitHub registry is used to store the docker image
4. Docker is used to build the flask image
5. AWS is the cloud provider
6. “aws-vault” tool (<https://github.com/99designs/aws-vault>) to access the EKS cluster via CLI
7. Lens(<https://k8slens.dev/>) and kubectl CLI are used to access the EKS cluster for troubleshooting

EKS cluster creation

The EKS cluster was created using Terraform.

1. Due to the security concern, an IAM user with permission to assume an IAM role will be used to deploy the cluster.
2. VPC, subnets and other network resources were created via the “vpc” module
3. EKS cluster and load balancer controller roles were created via the “eks” module

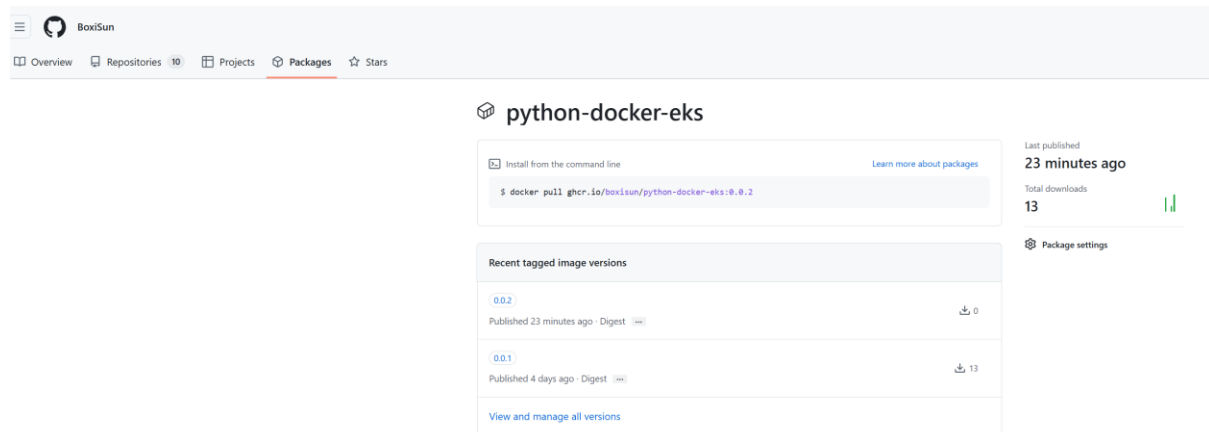
```
Apply complete! Resources: 60 added, 0 changed, 0 destroyed.

Outputs:

cluster_name = "flask-eks-cluster"
oidc_provider = "oidc.eks.eu-west-1.amazonaws.com/id/64E6DA64567D6249146380B45930EDBF"
oidc_provider_arn = "arn:aws:iam::972156694227:oidc-provider/oidc.eks.eu-west-1.amazonaws.com/id/64E6DA64567D6249146380B45930EDBF"
vpc_arn = "arn:aws:ec2:eu-west-1:972156694227:vpc/vpc-06cbaceb1aa6c4eb1"
vpc_cidr_block = "10.111.160.0/20"
vpc_id = "vpc-06cbaceb1aa6c4eb1"
```

Python flask application build and deployment

The flask application was built and pushed to GitHub as a package. A self-signed certificate was generated for the application server. (I roll back in my test with the AWS certificate as different certificates used by ALB and flask server will cause 502 Bad Gateway)



The flask application was deployed to the EKS cluster using Terraform.

```
Apply complete! Resources: 10 added, 0 changed, 0 destroyed.
```

Outputs:

```
ingress = "flask-ingress"
namespace = "flask"
rbac_role = "flask_rbac_role"
service = "flask-svc"
service_account = "flask-sa"
```

```
C:\Users\bsun\Projects\migrated\python-docker-eks\flask-app-deployment>kubectl get pods -n flask
NAME                                READY   STATUS    RESTARTS   AGE
flask-deployment-7675c66c78-5kqjh   1/1     Running   0           10m
flask-deployment-7675c66c78-clc4m   1/1     Running   0           10m
```

A TLS certificate was generated in AWS Certificate Manager to enable HTTPS connection.

A DNS CNAME was created in the company hosted zone in another AWS account via Terraform

```
[root@ip-10-111-163-19 ec2-user]#  
[root@ip-10-111-163-19 ec2-user]# curl https://flask.neon.markets  
{"message: "Hello, world!"}[root@ip-10-111-163-19 ec2-user]# ^C  
[root@ip-10-111-163-19 ec2-user]# curl http://flask.neon.markets  
<html>  
<head><title>301 Moved Permanently</title></head>  
<body>  
<center><h1>301 Moved Permanently</h1></center>  
</body>  
</html>  
[root@ip-10-111-163-19 ec2-user]# curl https://flask.neon.markets  
{"message: "Hello, world!"}[root@ip-10-111-163-19 ec2-user]#
```

Metrics server and horizontal pod autoscaling

The metrics server was deployed via “kubectl”

```
C:\Users\bsun\Projects\migrated\python-docker-eks\flask-app-deployment>kubectl apply -f https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml
serviceaccount/metrics-server created
clusterrole.rbac.authorization.k8s.io/system:aggregated-metrics-reader created
clusterrole.rbac.authorization.k8s.io/system:metrics-server created
rolebinding.rbac.authorization.k8s.io/metrics-server-auth-reader created
clusterrolebinding.rbac.authorization.k8s.io/metrics-server:system:auth-delegator created
clusterrolebinding.rbac.authorization.k8s.io/system:metrics-server created
service/metrics-server created
deployment.apps/metrics-server created
apiservice.apiregistration.k8s.io/v1beta1.metrics.k8s.io created

C:\Users\bsun\Projects\migrated\python-docker-eks\flask-app-deployment>kubectl get deployment metrics-server -n kube-system
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
metrics-server 1/1      1             1           91s

C:\Users\bsun\Projects\migrated\python-docker-eks\flask-app-deployment>kubectl top nodes
NAME          CPU(cores)   CPU%   MEMORY(bytes)   MEMORY%
ip-10-111-161-13.eu-west-1.compute.internal  38m      0%    627Mi           4%
ip-10-111-161-19.eu-west-1.compute.internal  29m      0%    657Mi           4%
ip-10-111-162-123.eu-west-1.compute.internal 34m      0%    627Mi           4%
ip-10-111-162-13.eu-west-1.compute.internal  26m      0%    630Mi           4%
```

```
C:\Users\bsun\Projects\migrated\python-docker-eks\flask-app-deployment>kubectl top pods -n flask
NAME          CPU(cores)   MEMORY(bytes)
flask-deployment-7675c66c78-5kqjh  2m           19Mi
flask-deployment-7675c66c78-clc4m  1m           18Mi
```

I run “while sleep 0.01; do wget -q -O- https://flask.neon.markets; done” in two sessions to trigger the horizontal autoscaling and monitor the resource usage

```
C:\Users\bsun\Projects>kubectl autoscale deployment flask-deployment -n flask --cpu-percent=2 --min=1 --max=10
horizontalpodautoscaler.autoscaling/flask-deployment autoscaled
```

Before scaling up:

```
C:\Users\bsun\Projects\migrated\sharedservices-platform>kubectl get hpa -n flask
NAME          REFERENCE          TARGETS   MINPODS   MAXPODS   REPLICAS   AGE
flask-deployment  Deployment/flask-deployment  0%/2%      1         10         1          88m

C:\Users\bsun\Projects\migrated\sharedservices-platform>kubectl top pods -n flask
NAME          CPU(cores)   MEMORY(bytes)
flask-deployment-7675c66c78-clc4m  1m           19Mi
```

After scaling up:

```
C:\Users\bsun\Projects\migrated\sharedservices-platform>kubect1 get hpa -n flask
NAME REFERENCE TARGETS MINPODS MAXPODS REPLICAS AGE
flask-deployment Deployment/flask-deployment 7%/2% 1 10 3 92m

C:\Users\bsun\Projects\migrated\sharedservices-platform>kubect1 get hpa -n flask
NAME REFERENCE TARGETS MINPODS MAXPODS REPLICAS AGE
flask-deployment Deployment/flask-deployment 7%/2% 1 10 3 92m

C:\Users\bsun\Projects\migrated\sharedservices-platform>kubectel top pods -n flask
'kubectel' is not recognized as an internal or external command,
operable program or batch file.

C:\Users\bsun\Projects\migrated\sharedservices-platform>kubect1 get hpa -n flask
NAME REFERENCE TARGETS MINPODS MAXPODS REPLICAS AGE
flask-deployment Deployment/flask-deployment 4%/2% 1 10 3 93m

C:\Users\bsun\Projects\migrated\sharedservices-platform>kubect1 get hpa -n flask
NAME REFERENCE TARGETS MINPODS MAXPODS REPLICAS AGE
flask-deployment Deployment/flask-deployment 3%/2% 1 10 5 93m

C:\Users\bsun\Projects\migrated\sharedservices-platform>kubect1 get hpa -n flask
NAME REFERENCE TARGETS MINPODS MAXPODS REPLICAS AGE
flask-deployment Deployment/flask-deployment 2%/2% 1 10 5 93m

C:\Users\bsun\Projects\migrated\sharedservices-platform>kubect1 top pods -n flask
NAME CPU(cores) MEMORY(bytes)
flask-deployment-7675c66c78-7ptw9 1m 18Mi
flask-deployment-7675c66c78-bfx9m 1m 18Mi
flask-deployment-7675c66c78-cjhg2 1m 18Mi
flask-deployment-7675c66c78-clc4m 21m 19Mi
flask-deployment-7675c66c78-q622s 1m 18Mi
```

```
C:\Users\bsun\Projects\migrated\sharedservices-platform>kubect1 top pods -n flask
NAME CPU(cores) MEMORY(bytes)
flask-deployment-7675c66c78-7ptw9 1m 18Mi
flask-deployment-7675c66c78-bfx9m 1m 18Mi
flask-deployment-7675c66c78-cjhg2 1m 18Mi
flask-deployment-7675c66c78-clc4m 20m 19Mi
flask-deployment-7675c66c78-q622s 1m 18Mi

C:\Users\bsun\Projects\migrated\sharedservices-platform>kubect1 top pods -n flask
NAME CPU(cores) MEMORY(bytes)
flask-deployment-7675c66c78-7ptw9 1m 18Mi
flask-deployment-7675c66c78-bfx9m 1m 18Mi
flask-deployment-7675c66c78-cjhg2 1m 18Mi
flask-deployment-7675c66c78-clc4m 20m 19Mi
flask-deployment-7675c66c78-q622s 1m 18Mi

C:\Users\bsun\Projects\migrated\sharedservices-platform>kubect1 top pods -n flask
NAME CPU(cores) MEMORY(bytes)
flask-deployment-7675c66c78-7ptw9 1m 18Mi
flask-deployment-7675c66c78-bfx9m 1m 18Mi
flask-deployment-7675c66c78-cjhg2 1m 18Mi
flask-deployment-7675c66c78-clc4m 21m 19Mi
flask-deployment-7675c66c78-q622s 1m 18Mi

C:\Users\bsun\Projects\migrated\sharedservices-platform>kubect1 top pods -n flask
NAME CPU(cores) MEMORY(bytes)
flask-deployment-7675c66c78-7ptw9 1m 18Mi
flask-deployment-7675c66c78-bfx9m 1m 18Mi
flask-deployment-7675c66c78-cjhg2 1m 18Mi
flask-deployment-7675c66c78-clc4m 21m 19Mi
flask-deployment-7675c66c78-q622s 1m 18Mi

C:\Users\bsun\Projects\migrated\sharedservices-platform>kubect1 top pods -n flask
NAME CPU(cores) MEMORY(bytes)
flask-deployment-7675c66c78-7ptw9 1m 18Mi
flask-deployment-7675c66c78-bfx9m 1m 18Mi
flask-deployment-7675c66c78-cjhg2 1m 18Mi
flask-deployment-7675c66c78-clc4m 21m 19Mi
flask-deployment-7675c66c78-q622s 1m 18Mi
```

With an php-apache server however, I could see that the workload are evenly distributed

kubectl apply -f <https://k8s.io/examples/application/php-apache.yaml>

kubectl run -i --tty load-generator --rm --image=busybox --restart=Never -- /bin/sh -c "while sleep 0.01; do wget -q -O- http://php-apache; done" (this is to trigger the scaling)

```
C:\Users\bsun>kubectl top pods
NAME                                CPU(cores)   MEMORY(bytes)
load-generator                      17m          0Mi
php-apache-598b474864-t5dv2        213m         11Mi

C:\Users\bsun>kubectl get hpa php-apache
NAME           REFERENCE             TARGETS   MINPODS   MAXPODS   REPLICAS   AGE
php-apache     Deployment/php-apache  106%/50%   1         10        1          60s

C:\Users\bsun>kubectl top pods
NAME                                CPU(cores)   MEMORY(bytes)
load-generator                      17m          0Mi
php-apache-598b474864-2ssm6        225m         11Mi
php-apache-598b474864-5q4gd        184m         11Mi
php-apache-598b474864-t5dv2        314m         11Mi
```

```
C:\Users\bsun>kubectl top pods
NAME                                CPU(cores)   MEMORY(bytes)
load-generator                      19m          0Mi
php-apache-598b474864-2ssm6        182m         11Mi
php-apache-598b474864-5q4gd        191m         11Mi
php-apache-598b474864-8wrfm        118m         11Mi
php-apache-598b474864-fd525        116m         11Mi
php-apache-598b474864-t5dv2        314m         11Mi

C:\Users\bsun>kubectl top pods
NAME                                CPU(cores)   MEMORY(bytes)
load-generator                      19m          0Mi
php-apache-598b474864-2ssm6        182m         11Mi
php-apache-598b474864-5q4gd        191m         11Mi
php-apache-598b474864-8wrfm        118m         11Mi
php-apache-598b474864-fd525        116m         11Mi
php-apache-598b474864-t5dv2        314m         11Mi

C:\Users\bsun>kubectl top pods
NAME                                CPU(cores)   MEMORY(bytes)
load-generator                      19m          0Mi
php-apache-598b474864-2ssm6        182m         11Mi
php-apache-598b474864-5q4gd        191m         11Mi
php-apache-598b474864-8wrfm        118m         11Mi
php-apache-598b474864-fd525        116m         11Mi
php-apache-598b474864-t5dv2        314m         11Mi
```

CICD with Bitbucket pipeline

Note: The Bitbucket pipeline requires the Bitbucket environment

The build step is triggered when there is a file change in the folder below:

filter:

includes:

- "flask-app-deployment/*"

Bug: Pulling the image of the same version will not overwrite the existing one.

The screenshot shows the Bitbucket Pipelines interface for a pipeline named 'marex-prod'. The pipeline consists of two steps: 'Build the flask docker image' and 'Deploy the flask image to the EKS cluster'. Both steps are marked as successful. The 'Build' step is expanded, showing the following commands and output:

```
Build docker

cd flask-docker

docker build --tag ghcr.io/boxisun/python-docker-eks:$version .

docker login --username BoxiSun --password $personal_token ghcr.io

docker push ghcr.io/boxisun/python-docker-eks:$version
```

The output of the 'docker push' command shows the image being pushed to the repository, with a digest of sha256:5244fd7e4a21a2b5cb7527242cc55894f5c9ea602f717fc31f5141518f65d225 and a size of 2283.

The deploy step:

The screenshot shows the Bitbucket Pipelines interface for a pipeline named 'marex-prod'. The pipeline consists of two steps: 'Build the flask docker image' and 'Deploy the flask image to the EKS cluster'. Both steps are marked as successful. The 'Deploy' step is expanded, showing the following commands and output:

```
Build docker

# kubernetes_deployment_v1.flask_deployment will be updated in-place
~ resource "kubernetes_deployment_v1" "flask_deployment" {
  id      = "flask/flask-deployment"
  # (1 unchanged attribute hidden)

  ~ spec {
    # (5 unchanged attributes hidden)

    ~ template {
      ~ spec {
        # (13 unchanged attributes hidden)

        ~ container {
          ~ image = "ghcr.io/boxisun/python-docker-eks:0.0.1" -> "ghcr.io/boxisun/python-docker-eks:0.0.4"
          name    = "flask"
          # (8 unchanged attributes hidden)

          # (3 unchanged blocks hidden)
        }
      }
    }
  }
  # (1 unchanged block hidden)

  # (2 unchanged blocks hidden)
}
# (1 unchanged block hidden)

# kubernetes_ingress_v1.flask_ingress will be updated in-place
~ resource "kubernetes_ingress_v1" "flask_ingress" {
  id      = "flask/flask-ingress"
  # (2 unchanged attributes hidden)

  ~ metadata {
    ~ annotations = {

```

After deployment:

```
[root@ip-10-111-163-19 ec2-user]# curl https://flask.neon.markets
{"message: "Hello, world!"}[root@ip-10-111-163-19 ec2-user]# curl https://flask.neon.markets
{"message: "Hello, world, test by Boxi!"}[root@ip-10-111-163-19 ec2-user]#
```


Argo CD

Argo CD is deployed to the EKS cluster by following the steps below:

1. Argo CD only works with yaml manifest files. I need to translate the terraform deployment code in the main.tf. The easiest way to do that is to run below command to export the flask deployment:

```
kubectl get deploy -n flask -o yaml > flask-app-deployment/app_manifests/flask-deployment.yaml
```

Note: The file exported contains some additional status information, I need to do some cleanup so that the file can be deployed. Try running “**kubectl apply -f flask-deployment.yaml**” until you get “deployment.apps/flask-deployment configured” without error.

2. Install argocd to the EKS cluster:

```
kubectl create namespace argocd
```

```
kubectl apply -n argocd -f https://raw.githubusercontent.com/argoproj/argo-cd/stable/manifests/install.yaml
```

```
C:\Users\bsun\Projects\migrated\python-docker-eks\flask-app-deployment>kubectl get pod -n argocd
NAME                                READY   STATUS    RESTARTS   AGE
argocd-application-controller-0     1/1     Running   0           53s
argocd-applicationset-controller-5ccbf46fd8-dqmjl  1/1     Running   0           54s
argocd-dex-server-7b6b9b5cb7-kl8x8  1/1     Running   0           54s
argocd-notifications-controller-5b9566d4c7-ppbsp  1/1     Running   0           54s
argocd-redis-665dd47d9b-fr76d       1/1     Running   0           54s
argocd-repo-server-fc97bfb4d-zdfpt   1/1     Running   0           53s
argocd-server-69754c4fb-vv574       1/1     Running   0           53s
```

3. Start Argo CD on your local computer and get the passport of “admin” to login the console:

```
kubectl port-forward -n argocd svc/argocd-server 8080:443
```

```
kubectl -n argocd get secret argocd-initial-admin-secret -o yaml
```

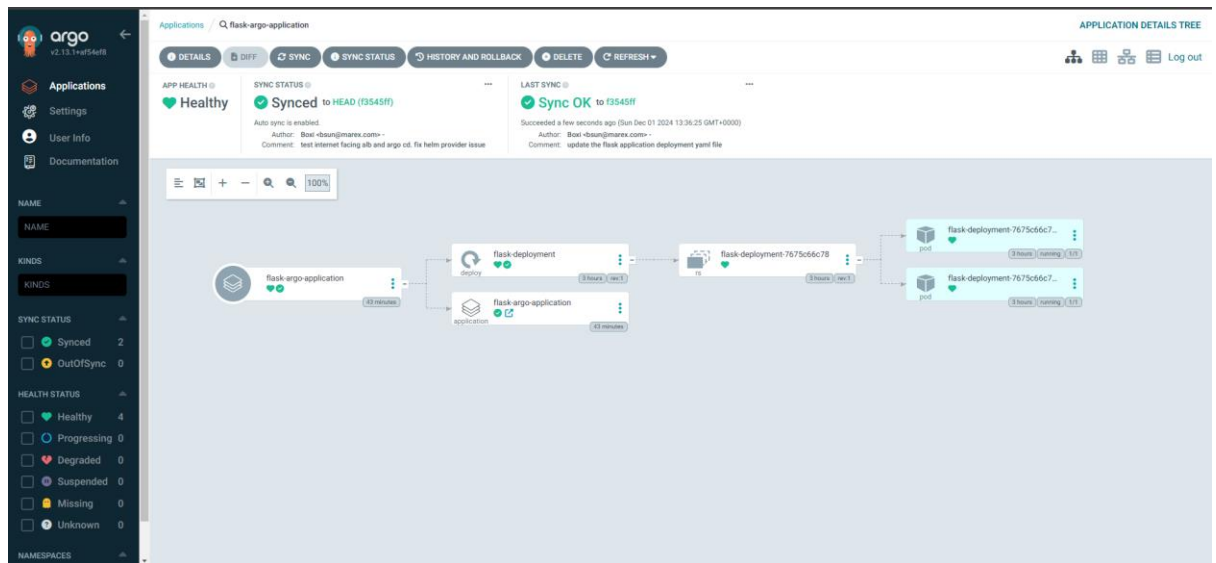
```
C:\Users\bsun\Projects\migrated\python-docker-eks\flask-app-deployment>kubectl -n argocd get secret argocd-initial-admin-secret -o yaml
apiVersion: v1
data:
  password: ZkJJTEVQM3oxZDlwOXVzLQ==
kind: Secret
metadata:
  creationTimestamp: "2024-12-01T12:16:31Z"
  name: argocd-initial-admin-secret
  namespace: argocd
  resourceVersion: "29917"
  uid: 221f2642-4a38-4799-8248-29466c5463e7
```

```
echo ZkJJTEVQM3oxZDlwOXVzLQ== | base64 --decode
```

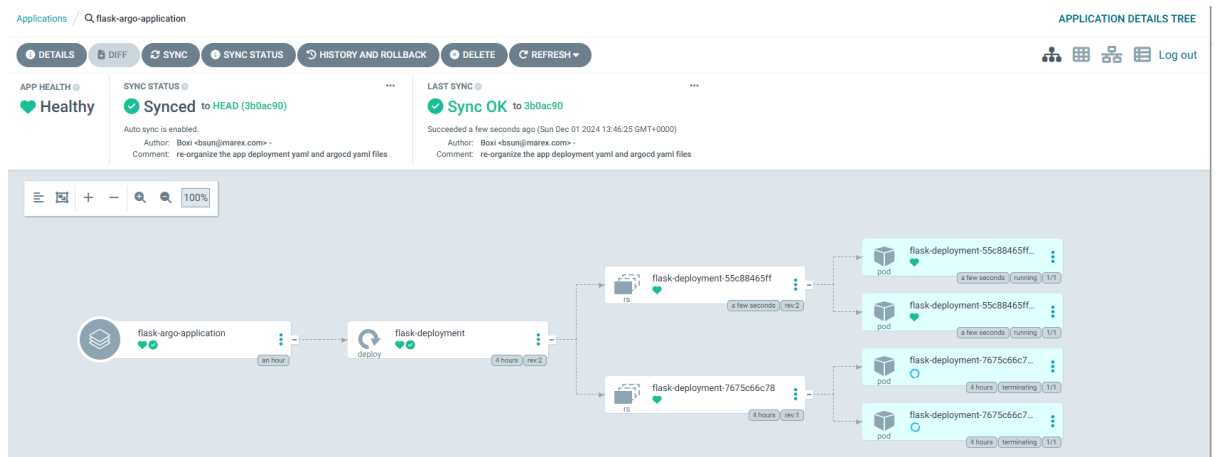
```
BSun@LDNSL4IT10 MINGW64 ~
$ echo ZkJJTEVQM3oxZDlwOXVzLQ== | base64 --decode
fBIEP3zld9p9us-
```

4. Link the flask GitHub repository to argocd:

```
kubectl apply -f flask-app-deployment/argocd/flask-argo-application.yaml
```



- Update the image version in the **flask-deployment.yaml** file and verify the change (tested twice with version 0.0.4 and 0.0.3):



https://flask.neon.markets

SaveShare

GEThttps://flask.neon.markets

Send

ParamsAuthorizationHeaders (7)BodyScriptsTestsSettings

Query Params

	Key	Value	Description	Bulk Edit
	Key	Value	Description	

BodyCookiesHeaders (5)Test Results

200 OK172 ms219 BSave Response

PrettyRawPreviewVisualizeHTML

1{"message": "Hello, world, test by Boxi!"}

DETAILSDIFFSYNCSYNC STATUSHISTORY AND ROLLBACKDELETEREFRESH

APP HEALTH

Healthy

SYNC STATUS

Synced to HEAD (3b5bd5f)

Auto sync is enabled.
Author: Boxi <bsun@marex.com>
Comment: re-organize the app deployment yaml and argo yaml files

LAST SYNC

Sync OK to 3b5bd5f

Succeeded a few seconds ago (Sun Dec 01 2024 14:19:06 GMT+0000)
Author: Boxi <bsun@marex.com>
Comment: test again by updating the image version to 0.0.3

GEThttps://flask.neon.markets

Send

ParamsAuthorizationHeaders (7)BodyScriptsTestsSettings

Query Params

	Key	Value	Description	Bulk Edit
	Key	Value	Description	

BodyCookiesHeaders (5)Test Results

200 OK180 ms226 BSave Response

PrettyRawPreviewVisualizeHTML

1{"message": "Hello, world, test by Boxi, again!"}

Issue faced and solutions applied

1. Can't connect to the API group after creating the EKS cluster

Error:

```
C:\Users\bsun\Projects>kubectl get svc
E1124 17:59:53.250966 25816 memcache.go:265] couldn't get current server API
group list: Get "https://90F2F70114A39F575C0197231A604818.gr7.eu-west-
1.eks.amazonaws.com/api?timeout=32s": dial tcp 10.111.162.76:443: i/o timeout
E1124 18:00:23.260130 25816 memcache.go:265] couldn't get current server API
group list: Get "https://90F2F70114A39F575C0197231A604818.gr7.eu-west-
1.eks.amazonaws.com/api?timeout=32s": dial tcp 10.111.162.76:443: i/o timeout
```

Troubleshooting:

- Check the firewall (AWS security groups) and make sure they are not blocking my IP.
- Since the EKS cluster is deployed in private subnets, the endpoint access is set to “private” meaning that it could only be accessible from within the VPC.

Solutions:

Create an EC2 jumpbox in the same private subnet of the EKS cluster and confirmed that I could access the cluster

Per the document below, change the endpoint access from ”private” to “Public and private”.

The screenshot shows the AWS Management Console interface for the 'flask-eks-cluster'. The breadcrumb navigation at the top reads: EKS > Clusters > flask-eks-cluster > Networking > Manage endpoint access. On the left sidebar, the 'Amazon Elastic Kubernetes Service' section is expanded, showing 'Clusters' and 'Related services' (Amazon ECR, AWS Batch). A blue banner at the top of the main content area states: 'Cluster termination in progress flask-eks-cluster is now being terminated. This process may take several minutes.' Below this, the heading 'Manage endpoint access: flask-eks-cluster' is displayed. The 'Cluster endpoint access' section includes an 'Info' icon and the text 'Configure access to the Kubernetes API server endpoint.' There are three radio button options: 'Public' (unselected), 'Public and private' (selected), and 'Private' (unselected). Each option has a descriptive text block explaining the traffic flow. At the bottom of this section is a link for 'Advanced settings'.

[EKS](#) > [Clusters](#) > [flask-eks-cluster](#) > [Networking](#) > Manage endpoint access

Amazon Elastic Kubernetes Service

[Clusters](#)

Amazon EKS Anywhere

Enterprise Subscriptions [New](#)

Related services

Amazon ECR

AWS Batch

Console settings

Documentation [🔗](#)

Submit feedback

Cluster termination in progress
flask-eks-cluster is now being terminated. This process may take several minutes.

Manage endpoint access: flask-eks-cluster

Cluster endpoint access [Info](#)
Configure access to the Kubernetes API server endpoint.

☐ Public
The cluster endpoint is accessible from outside of your VPC. Worker node traffic will leave your VPC to connect to the endpoint.

☒ **Public and private**
The cluster endpoint is accessible from outside of your VPC. Worker node traffic to the endpoint will stay within your VPC.

☐ Private
The cluster endpoint is only accessible through your VPC. Worker node traffic to the endpoint will stay within your VPC.

► **Advanced settings**

2. Load balancer is not created and terraform creating ingress timeout

Error:

Error: Load Balancer is not ready yet

```
with kubernetes_ingress_v1.sample_application_ingress,  
on sample_app.tf line 175, in resource "kubernetes_ingress_v1" "sample_application_ingress":  
175: resource "kubernetes_ingress_v1" "sample_application_ingress" {
```

Troubleshooting:

kubectl describe ingress sample-application-ingress -n sample-application

```
Message  
-----  
Failed build model due to couldn't auto-discover subnets: unable to resolve at least one subnet (0 match VPC and tags: {kubernetes.io/role/elb})
```

Solutions:

The error indicates that the alb couldn't find any subnet match. Per the document below, I choose to specify subnets as annotations

<https://docs.aws.amazon.com/eks/latest/userguide/alb-ingress.html>

- Your public and private subnets must meet the following requirements. This is unless you explicitly specify subnet IDs as an annotation on a service or ingress object. Assume that you provision load balancers by explicitly specifying subnet IDs as an annotation on a service or ingress object. In this situation, Kubernetes and the AWS load balancer controller use those subnets directly to create the load balancer and the following tags aren't required.
 - **Private subnets** – Must be tagged in the following format. This is so that Kubernetes and the AWS load balancer controller know that the subnets can be used for internal load balancers. If you use eksctl or an Amazon EKS AWS CloudFormation template to create your VPC after March 26, 2020, the subnets are tagged appropriately when created. For more information about the Amazon EKS AWS CloudFormation VPC templates, see [Create an Amazon VPC for your Amazon EKS cluster](#).
 - **Key** – `kubernetes.io/role/internal-elb`
 - **Value** – `1`
 - **Public subnets** – Must be tagged in the following format. This is so that Kubernetes knows to use only the subnets that were specified for external load balancers. This way, Kubernetes doesn't choose a public subnet in each Availability Zone (lexicographically based on their subnet ID). If you use eksctl or an Amazon EKS AWS CloudFormation template to create your VPC after March 26, 2020, the subnets are tagged appropriately when created. For more information about the Amazon EKS AWS CloudFormation VPC templates, see [Create an Amazon VPC for your Amazon EKS cluster](#).
 - **Key** – `kubernetes.io/role/elb`
 - **Value** – `1`

<https://kubernetes-sigs.github.io/aws-load-balancer-controller/v2.2/guide/ingress/annotations/>

- `alb.ingress.kubernetes.io/subnets` specifies the [Availability Zone](#) that ALB will route traffic to. See [Load Balancer subnets](#) for more details.

You must specify at least two `subnets` in different AZ. both `subnetID` or `subnetName`(Name tag on `subnets`) can be used.

Tip

You can enable `subnet` auto discovery to avoid specify this annotation on every Ingress. See [Subnet Discovery](#) for instructions.

Example

```
alb.ingress.kubernetes.io/subnets: subnet-xxxx, mySubnet
```



3. EKS failed to pull the docker image from github

Error:

Pulling image "ghcr.io/boxisun/python-docker-eks:0.0.1"

Source	kubelet ip-10-215-144-9.eu-west-1.compute.internal
Count	4
Sub-object	spec.containers{python-flask}
Last seen	2024-11-22T16:44:36Z

Failed to pull image "ghcr.io/boxisun/python-docker-eks:0.0.1": failed to pull and unpack image "ghcr.io/boxisun/python-docker-eks:0.0.1": failed to resolve reference "ghcr.io/boxisun/python-docker-eks:0.0.1": failed to authorize: failed to fetch anonymous token: unexpected status from GET request to https://ghcr.io/token?scope=repository%3Aboxisun%2Fpython-docker-eks%3Apull&service=ghcr.io: 401 Unauthorized

Source kubelet ip-10-215-144-9.eu-west-1.compute.internal

Troubleshooting:

The error indicates that EKS failed to fetch the token and therefore not authorized to pull the image

Solution:

1. Change the flask image to be public:

Manage access

1 member

BoxiSun

Change package visibility

☒ **Public**
Make this package visible to anyone.

☐ **Private**
Make this package visible privately, to organization members that have access.

Please type python-docker-eks to confirm: *

python-docker-eks

[I understand the consequences, change package visibility](#)

Danger Zone

Change package visibility
This package is currently private.

[Change visibility](#)

Delete this package
Once you delete a package, there is no going back. Please be certain.

[Delete this package](#)

2. Create a personal access token and follow the steps in the document below to generate a dockerconfigjson secret. Apply the secret and add the “image_pull_secret” to the deployment terraform code:

```
resource "kubernetes_deployment_v1" "flask_deployment" {  
  spec {  
    replicas = 2  
  
    selector {  
      match_labels = {  
        app = "${var.app_name}"  
      }  
    }  
  
    template {  
      metadata {  
        labels = {  
          app = "${var.app_name}"  
        }  
      }  
      spec {  
        service_account_name = kubernetes_service_account_v1.service-account-flask.metadata[0].name  
  
        #NOTE:image_pull_secrets is for private images. It requires the eks secrets to be created. The manifest can be found in the folder "pu  
        # image_pull_secrets {  
        #   name = "${var.app_name}-image-pull-credential"  
        # }  
      }  
    }  
  }  
}
```


4. Failed with “502 Bad Gateway” when testing the endpoint:

Error:

GET http://internal-k8s-sampleapp-sampleapp-fc69043995-1764473567.eu-west-1.elb.amazonaws.com/ Send

Params Authorization Headers (8) Body Scripts Tests Settings Cookies

Query Params

Key	Value	Description
Key	Value	Description

Body Cookies Headers (5) Test Results 502 Bad Gateway 41 ms 277 B Save Response

Pretty Raw Preview Visualize HTML

```
1 <html>
2
3 <head>
4   <title>502 Bad Gateway</title>
5 </head>
6
7 <body>
8   <center>
9     <h1>502 Bad Gateway</h1>
10  </center>
11 </body>
12
13 </html>
```

Troubleshooting:

Check the status of the load balancer target group and find out that all instances are “unhealthy”. This means the health check failed.

4 Total targets

0 Healthy 4 Unhealthy 0 Unused 0 Initial 0 Draining

0 Anomalous

▼ Distribution of targets by Availability Zone (AZ)

Select values in this table to see corresponding filters applied to the Registered targets table below.

Last fetched 16 minutes ago

	Total targets	Healthy	Unhealthy	Unused	Initial
euw1-az2	2	0	2	0	0
euw1-az1	2	0	2	0	0

Targets Monitoring Health checks Attributes Tags

Health check settings Edit

Protocol HTTP	Path /	Port Traffic port	Healthy threshold 2 consecutive health check successes
Unhealthy threshold 2 consecutive health check failures	Timeout 5 seconds	Interval 15 seconds	Success codes 200

In this case, the “Traffic port” of flask application is 5000. The deployment log confirms that:

```
kubectl logs -n sample-application sample-application-deployment-79c75f45b8-9rqv5
```

* Debug mode: off

WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.

* Running on all addresses (0.0.0.0)

* Running on http://127.0.0.1:5000

* Running on http://100.64.0.241:5000

Press CTRL+C to quit

Solutions:

I found from the terraform code that the container port and target_port were both 80. I updated them to 5000 and the issue was fixed

```
spec {
  service_account_name = kubernetes_service_account.service-account-sample-app.metadata[0].name
  container {
    image = "ghcr.io/boxisun/python-docker-eks:0.0.1"
    name  = "python-flask"

    resources {
      limits = {
        cpu    = "0.5"
        memory = "512Mi"
      }
      requests = {
        cpu    = "250m"
        memory = "50Mi"
      }
    }
  }

  port {
    container_port = 5000
  }
}
```

```
resource "kubernetes_service_v1" "sample_application_svc" {
  metadata {
    name      = "sample-application-svc"
    namespace = kubernetes_namespace.sample-application-namespace.metadata[0].name
  }
  spec {
    selector = {
      app = "nginx"
    }
    session_affinity = "ClientIP"
    port {
      port      = 80
      target_port = 5000
    }

    type = "NodePort"
  }
}
```

5. Pods reboot over and over again

Error:

```
Liveness probe failed: Get "http://10.111.161.34:80/": dial tcp 10.111.161.34:80: connect: connection refused
```

Troubleshooting:

Similar to Issue 4, the port 80 in the error got my attention. Plus, the “initial_delay_seconds” and “period_seconds” setting are way too short.

Solutions:

I updated the terraform code and apply the change

```
~ liveness_probe {  
  ~ initial_delay_seconds = 3 -> 10  
  ~ period_seconds       = 3 -> 10  
  # (3 unchanged attributes hidden)  
  
  ~ http_get {  
    ~ port = "80" -> "5000"  
    # (2 unchanged attributes hidden)
```